

PodSupervisor

Are your pods sleeping?

Glenn West

gwest@redhat.com

Sep 30, 2022 10:31

Issue

- In a highly redundant OpenShift environment – Fail Fast is important to uptime of applications.
- Script “kills” pods on not-ready node and pod will be rescheduled as its settings dictate.




Solution

- A new monitor script that watches all nodes for not-ready then cleans up selected pods
- Pods may be selected list of names or wildcards
- Pods may be avoided by list of names or wildcard
- Script auto suspends during upgrades
- Default settings watch dns pods.

```
↑ ... @@ -5,12 +5,20 @@ input="/tmp/pods.data"
5 5 while IFS= read -r line
6 6 #Example data
7 7 #testapp httpd-ex-1-build 0/1
8 + #openshift-dns dns-default-kzgfg 3/3
9 9 do
10 10 #echo "$line"
11 11 namespace=`echo $line | cut -d' ' -f1`
12 12 podname=`echo $line | cut -d' ' -f2`
13 13 nodename=`echo $line | cut -d' ' -f8`
14 14 ocpstr='openshift'
15 + if [[ "$namespace" == "openshift-dns" ]]; then
16 +     if [[ "$nodename" == $1 ]]; then
17 +         echo "Cleaning up $nodename/$namespace/$podname"
18 +         echo "PodSupervisor: Node: $nodename Application: $namespace Pod: $podname - Deleted" > /dev/log
19 +         oc delete po/$podname --namespace $namespace
20 +     fi
21 + fi
14 22 if [[ "$namespace" != *$ocpstr* ]]; then
```

Additional Features

- Node flapping detection
 - Node must be not ready two scans
- On pod deletion, alerts are sent to OpenShift Alert Manager.



Example

- Node is marked
- Script is run
- Two pods are deleted

```
gwest@gwest@redhat podsupervisor % oc get nodes
NAME                                STATUS    ROLES    AGE   VERSION
control-plane-0                    Ready     master   27h   v1.20.10+bbbc079
control-plane-1                    Ready     master   27h   v1.20.10+bbbc079
control-plane-2                    Ready     master   27h   v1.20.10+bbbc079
worker-0                           Ready     worker   27h   v1.20.10+bbbc079
worker-1                           Ready,SchedulingDisabled worker   27h   v1.20.10+bbbc079
gwest@gwest@redhat podsupervisor % oc get pods
NAME                                READY    STATUS    RESTARTS   AGE
httpd-ex-7bd6c6788-7hlcj           1/1      Running   0           151m
httpd-ex-7bd6c6788-rdfdc           1/1      Running   0           151m
gwest@gwest@redhat podsupervisor % ./podsupervisor.sh
Cleaning up worker-1
Cleaning up worker-1/testapp/httpd-ex-7bd6c6788-7hlcj
pod "httpd-ex-7bd6c6788-7hlcj" deleted
Cleaning up worker-1/testapp/httpd-ex-7bd6c6788-rdfdc
pod "httpd-ex-7bd6c6788-rdfdc" deleted
gwest@gwest@redhat podsupervisor % █
```

Result

- Pods have moved to worker-0.
- No system/openshift pods were touched or harmed.

```
[gwest@gwest@redhat podsupervisor % oc get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
httpd-ex-7bd6c6788-7pvkt	1/1	Running	0	80s	10.131.0.32	worker-0	<none>	<none>
httpd-ex-7bd6c6788-lt269	1/1	Running	0	77s	10.131.0.33	worker-0	<none>	<none>

Source

- Upstream source:
- <https://github.com/glennswest/podsupervisor>